

**CS310 Natural Language Processing**  
**Assignment 5: Dependency Parsing**  
**Total points: 50 + (10 bonus)**

**Tasks**

- Train a feed-forward neural network-based dependency parser and evaluate its performance on the provided treebank dataset.
- (Bonus 1) Implement the [arc-eager approach](#) for the parser
- (Bonus 2) Re-implement the [feature extractor using Bi-LSTM as the backbone encoder](#).

**Submit**

- See REAME.md for submission requirements.
- Report document is not mandatory for Task 1-5. It is required for bonus Task 6 and 7.

**Requirements**

**1) Implement the Feature Extractor (15 points)**

Following Chen and Manning (2014), the parser will use features [from the top three words on stack](#)  $s_0, s_1, s_2$ , the top three words on buffer  $b_0, b_1, b_2$ . Using the implementation from the lab practice, they correspond to `stack[-1]`, `stack[-2]`, `stack[-3]`, and `buffer[-1]`, `buffer[-2]`, `buffer[-3]`, respectively. We also include their POS tags as feature, i.e.,  $t_1, \dots, t_n$ .

Specifically, words and tags are first associated with embedding vectors:  $e(w_i)$  and  $e(t_i)$  for  $i = 1, \dots, n$ , where  $n$  is the length of input sentence. Then the feature for the current configuration  $c$  is:

$$\phi(c) = e(s_2) \oplus e(s_1) \oplus e(s_0) \oplus e(b_0) \oplus e(b_1) \oplus e(b_2) \oplus e(ts_2) \oplus e(ts_1) \oplus e(ts_0) \oplus e(tb_0) \oplus e(tb_1) \oplus e(tb_2)$$

Here,  $ts_i$  is the POS tag of the  $i$ th word on stack, and  $tb_i$  is the tag of the  $i$ th word on buffer. Note that they are **NOT the  $i$ th word in the sentence**.

In some configurations, the stack or buffer may have fewer than 3 words. In those cases, use pseudo tokens “<NULL>” to replace the missing blanks, which is also associated with an embedding  $e("\langle NULL \rangle")$ . So is the special token “<ROOT>”,  $e("\langle ROOT \rangle")$ . [Their POS tags \(which do not exist\) can be some pseudo values as well.](#)

For example, if the buffer contains [“apple”, “trees”, “grow”] and the stack contains [“<ROOT>”, “the”], then the concatenated word vectors are:

$$e("\langle NULL \rangle") \oplus e("\langle ROOT \rangle") \oplus e("the") \oplus e("apple") \oplus e("trees") \oplus e("grow")$$

You can infer the concatenated tag vectors similarly.

Chen et. al. (2014) uses embedding size  $d = 50$ . You can use [larger values](#).

**2) Implement Scoring Oracle (10 points)**

The feature extracted will be passed to the scoring **oracle** that determines the transition action  $t$  given the feature  $\phi(c)$  extracted from the feature function. The scoring function could be multiple layer perceptron (MLP):  $Score_\theta(\phi(c), t) = MLP_\theta(\phi(c))[t]$ , where  $MLP_\theta(x) = W^{[2]} \cdot \tanh(W^{[1]} \cdot x + b^{[1]}) + b^{[2]}$

(Chen et. al. (2014) uses hidden layer size  $h = 200$  and tanh activation. You can use ReLU.)

The oracle model should be implemented [as a sub-class of `torch.nn.Module`](#). Two versions of oracle are required: `BaseModel` (using words only) and `WordPOSModel` (using words + POS)

### 3) Training (5 points)

In which the `forward` function takes input a sequence of training data instances A training instance has two parts,  $X$  and  $y$ .  $X$  represents the current configuration states, that is, the integer IDs of the 6 words in stack and buffer, and their POS tag IDs. You should implement how these IDs are converted to embedding vectors as described in Task 1 and 2, and then compute the loss(negative log-likelihood loss) between the predicted transition action  $\hat{y}$  and the ground truth  $y$ .

### 4) Implement the Parser for Inference (15 points)

We want the parser to be able to parse an unannotated input sentence, so you should also implement a `parse_sentence` function (as a member of a `Parser` class) that takes a sequence of words as input, and return the parsed tree. This function is very similar to `forward in terms of computation`, with the exception that it does not know the gold-standard transition actions. This function's job is to accomplish the following transition-based greedy parsing:

---

**Algorithm 1** Greedy transition-based parsing

---

```
1: Input: sentence  $s = w_1, \dots, x_w, t_1, \dots, t_n$ ,  
   parameterized function  $\text{SCORE}_\theta(\cdot)$  with param-  
   eters  $\theta$ .  
2:  $c \leftarrow \text{INITIAL}(s)$   
3: while not  $\text{TERMINAL}(c)$  do  
4:    $\hat{t} \leftarrow \arg \max_{t \in \text{LEGAL}(c)} \text{SCORE}_\theta(\phi(c), t)$   
5:    $c \leftarrow \hat{t}(c)$   
6: return  $\text{tree}(c)$ 
```

---

Figure from Kiperwasser and Goldberg (2016)

There are several notable places of the above algorithm:

- ✧ You can implement the exit condition  $\text{TERMINAL}(c)$  using an **if** statement to check **if “<ROOT>” is the only element on stack and the buffer is empty.**
- ✧ The argmax operation means that you select the highest scoring transition, but unfortunately it is possible that the highest scoring transition is not possible. Therefore, instead of selecting the highest-scoring action, you should select **the highest scoring permitted transition**, indicated by the  $t \in \text{LEGAL}(c)$  subscript.

### 5) Evaluation (5 points)

Run the evaluation script and check the LAS and UAS scores on the dev and test sets.

**\*\* Grading rubrics \*\***

- If your model is implemented correctly and the training code can run without problem, then you get the full credits for Task 1, 2, and 3.
- If your UAS score  $x \geq 70\%$  on the **test** set, then you get **full points** for Task 4 and 5.
- If  $50\% \leq x < 60\%$ , you get **10 points** for Task 4 and **2.5 points** for Task 5
- If  $x < 50\%$ , then you receive **0 points** for Task 4 and 5

**\*\*\*\***

### Bonus Tasks

#### 6) Implement the arc-eager approach (5 points)

You only need to modify the `get_training_instances` function and the `State` class. You do not need to actually train the parser. Note that you should add the new “Reduce” action to the code. Prove that your implementation is different from the arc-standard approach by showing an example. [\(The example should come in the report\)](#)

## 7) Implement the Bi-LSTM-based encoder (5 points)

Following Kiperwasser and Goldberg (2016), which uses words and POS tags as features. Given a  $n$ -words input sentence  $w_1, \dots, w_n$  together with the corresponding POS tags  $t_1, \dots, t_n$ , we associate each word and POS tag with an embedding vector  $e(w_i)$  and  $e(t_i)$ , and create a sequence of input vectors  $x_{1:n}$  in which each  $x_i$  is the concatenation of the word and POS embeddings:

$$x_i = e(w_i) \oplus e(t_i)$$

The input  $x_{1:n}$  is then passed to a Bi-LSTM encoder, which produces a hidden representation for each input element:

$$v_i = \text{BiLSTM}(x_i, i)$$

The feature function is the concatenated Bi-LSTM vectors of the top 3 items on the stack and the first item on the buffer. That is, for a configuration  $c = \{\text{stack: } [\dots | s_2 | s_1 | s_0], \text{buffer: } [b_0 | \dots]\}$ , the feature for this configuration is:

$$\phi(c) = v_{s_2} \oplus v_{s_1} \oplus v_{s_0} \oplus v_{b_0}, \text{ where } v_i = \text{BiLSTM}(x_{1:n}, i)$$

In our implementation for the configuration state in the lab, the stack top  $s_0$  corresponds to `state.stack[-1]`, and the buffer front  $b_0$  corresponds to `state.buffer[-1]`, and so forth for other elements. If the stack contains fewer than 3 words or the buffer is empty, then use the special token “<NULL>” to fill up the blanks.

For example, if the buffer = [..., “apple”] and stack = [<“<ROOT>”, “the”], then we should use  $e(\langle \text{NULL} \rangle) \oplus e(\langle \text{ROOT} \rangle) \oplus v_{\text{the}} \oplus v_{\text{apple}}$  as the feature.

For another example, if the buffer is empty and stack = [<“<ROOT>”, “the”], then we should use  $e(\langle \text{NULL} \rangle) \oplus e(\langle \text{ROOT} \rangle) \oplus v_{\text{the}} \oplus e(\langle \text{NULL} \rangle)$  as the feature

Note that here the vectors for “<NULL>” and “<ROOT>” are not the output from Bi-LSTM because they not actual words, and they are just returned by the embedding table.

[\(If your Bi-LSTM implementation is complex, explain your code and demonstrate how performance changes in the report\)](#)

### \*\* Grading rubrics \*\*

- For Task 6, you only need to change the `State` class and the `get_training_instances` function by adding “Reduce” to the transition action set. You do not need to train a new parser.
- For Task 7, you need to have better LAC and UAC scores than Task 5.

## References

Chen, D. and C. D. Manning (2014). [A fast and accurate dependency parser using neural networks](#). Proceedings of the 2014 conference on empirical methods in natural language processing (EMNLP).

Kiperwasser, E. and Y. Goldberg (2016). “Simple and accurate dependency parsing using bidirectional LSTM feature representations.” [Transactions of the Association for Computational Linguistics](#) 4: 313–327.